

# Agentic RAG

Agent Framework Dev Project  
Jan 2026 - Burlington MA edition

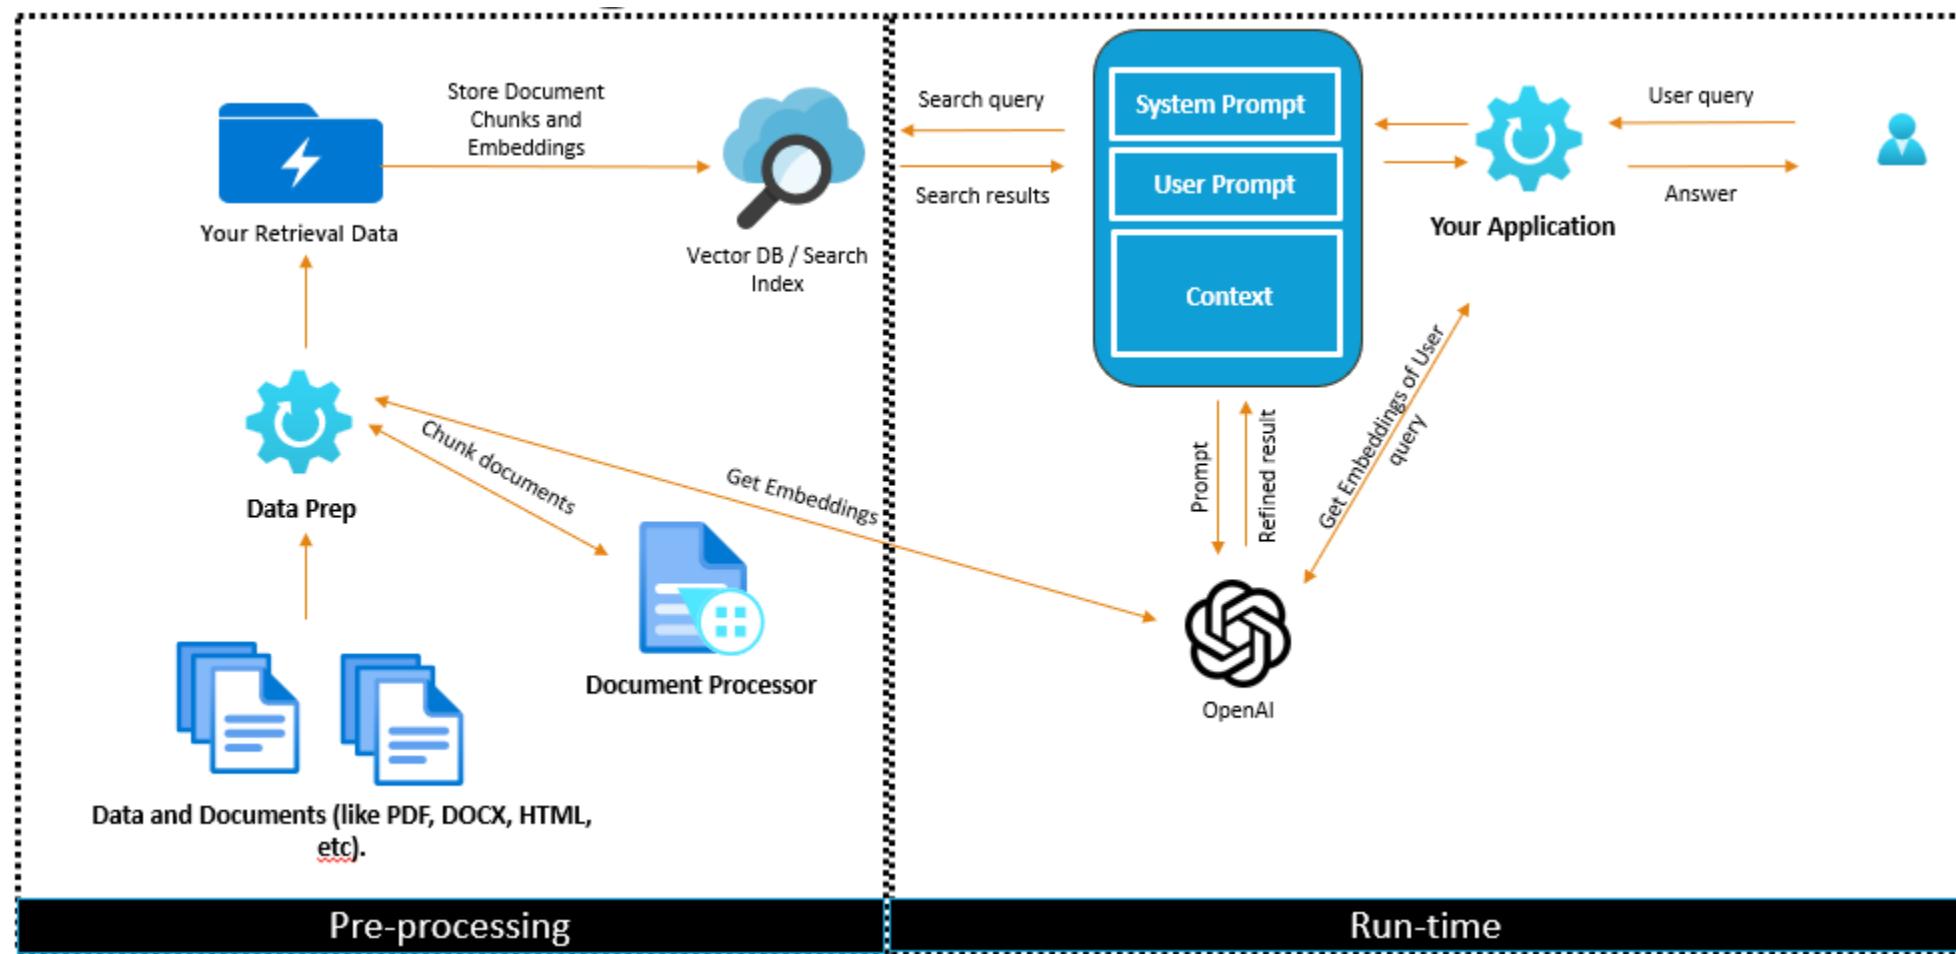
Jason Haley  
Microsoft AI MVP

jason@jasonhaley.com

# Today's Overview

- Quick recap of traditional RAG
- Why basic RAG breaks down
- What makes agentic RAG “agentic”
- Core concepts and architecture
- When to use agentic RAG
- Key takeaways

# Traditional RAG in 60 Seconds



# Traditional RAG: Quick Recap

## The Basic Flow

User query → Embed → Similarity search → Retrieve context → Generate answer

## The Problem It Solves

Grounds LLM responses in your data, reducing hallucinations and providing source attribution

## Simple & Effective

Works great for straightforward queries with clear intent

# Why Traditional RAG Breaks Down

## Example Query

*"What were the key findings from our Q3 reports and how do they compare to industry trends?"*

### Fixed Retrieval

Single similarity search can't break down complex queries

### Single Source

Can't search internal docs and external data

### No Iteration

One shot only, can't refine or follow up

## The Core Problem

Traditional RAG treats all queries the same way.

# What Makes Agentic RAG “Agentic”?

## The LLM becomes a reasoning engine

It decides what to do, which tools to use, and when to stop



### Planning

Break complex queries  
into steps



### Tool Selection

Choose the right data  
source



### Iteration

Refine until satisfied

# The Agents' Jobs



**Example:** "What's our pricing vs competitors?"

**Thought:** Need internal pricing first

**Action:** search\_internal\_docs("pricing")

**Observation:** Found our pricing → \$49/month

**Thought:** Now need competitor data

**Action:** web\_search("competitor pricing")

**Observation:** Found competitors → Average \$55/month

# Query Augmentation & Multi-Step Retrieval

## Example Query

"What were the key findings from our Q3 reports and how do they compare to industry trends?"

## Query Rewriting

Transform original user query to be effective for your retrieval system

## Query Expansion

Enhance results by generating multiple queries from the original query

## Query Decomposition

Breakdown complex user queries into focused sub-queries

# Retrieval Is Just One Tool

## Vector Store

Semantic Searches

## Web Search

Realtime External Data

## SQL/APIs

Structured Data Searches

## Code Interpreter

Tools Using Code to Analyze Data

## Document Diffing

Tools for Advanced Document Comparison

## Calculators

Tools to Aggregate or Perform Complex Math

# When to Use Agentic RAG

## Traditional RAG

### When to use:

- Simple, direct queries
- Single data source
- Speed is critical
- Low cost requirement
- Predictable behavior needed

### Example:

*"What is our return policy?"*

## Agentic RAG

### When to use:

- Complex, multi-part queries
- Multiple data sources
- Query decomposition needed
- Iterative refinement valuable
- Flexibility over speed

### Example:

*"Compare our Q3 performance to industry trends and suggest areas for improvement"*

# Key Takeaways

## **Agentic RAG = RAG + Reasoning**

The LLM plans, selects tools, and iterates until it has enough information

### **Best for complex, multi-source queries**

Use when query complexity justifies the additional cost and latency

### **Trade flexibility for predictability**

More intelligent but less deterministic than traditional RAG

### **Start small and iterate**

2-3 tools, clear descriptions, comprehensive logging